

A Framework for an In-depth Comparison of Scale-up and Scale-out Systems

Michael Sevilla

Ike Nassi, Kleoni Ioannidou,
Scott Brandt, Carlos Maltzahn

{msevilla, inassi, kleoni, scott, carlosm}@soe.ucsc.edu

University of California, Santa Cruz

December 19, 2013

Scaling

scale-up
vs.
scale-out

Michael Sevilla

Q: *What do we do when there is too much data?*


A: **Scale** the system

► out

- ++ nodes to the system
- modify applications

► up

- ++ resources to a single node
- modify the system



Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

Scaling


scale-up
vs.
scale-out

Michael Sevilla

Q: *What do we do when there is too much data?*

A: **Scale** the system

- ▶ out
 - ++ nodes to the system
 - modify applications
- ▶ up
 - ++ resources to a single node
 - modify the system



Q: *Which is better?*

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

Contributions

1. Comparison framework for scale-out/up
2. Achieve scale-out properties on scale-up
 - ▶ Parallelism limited by new job phases
 - ▶ Fault tolerance can make scale-up slower than scale-out
 - ▶ Scalable storage may be the ultimate bottleneck

We show:

- ▶ must consider properties when comparing scale-out/up
- ▶ limitations of a scale-up computation framework

scale-up
vs.
scale-out

Michael Sevilla

Introduction
contributions
challenges


Methodology
parameters
input
software
hardware

Analysis
initial results
implementations
parallelism


Conclusion

Goal

Framework for comparing:



vs.



Why re-examine scale-up?

- ▶ new technology
- ▶ simplicity
- ▶ legacy applications

scale-up
vs.
scale-out

Michael Sevilla

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware


Analysis

initial results
implementations
parallelism


Conclusion

Goal

Framework for comparing:



vs.



Limit study to MapReduce

- ▶ standard for big data analytics
- ▶ goal is to make fair comparisons

scale-up
vs.
scale-out

Michael Sevilla

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

Challenges - How do we:

Q: *compare algorithms?*

Q: *compare hardware?*

Q: *account for properties provided by scale-out by default?*

By design, scale-out provides:

- ▶ parallelism by automatically distributing load
- ▶ fault tolerance by rescheduling computation
- ▶ scalable storage with a distributed file system
- ▶ portability; Hadoop applications can run on any cluster
- ▶ availability because it can continually service clients
- ▶ scalability

scale-up
vs.
scale-out

Michael Sevilla

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

Challenges - How do we:

Q: *compare algorithms?*

Q: *compare hardware?*

Q: *account for properties provided by scale-out by default?*

By design, scale-out provides:

- ▶ parallelism by automatically distributing load
- ▶ fault tolerance by rescheduling computation
- ▶ scalable storage with a distributed file system
- ▶ portability; Hadoop applications can run on any cluster
- ▶ availability because it can continually service clients
- ▶ scalability

These properties may or may not affect performance...

...but they can't be ignored!

scale-up
vs.
scale-out

Michael Sevilla

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

Challenges - How do we:

Q: *compare algorithms?*

Q: *compare hardware?*

Q: *account for properties provided by scale-out by default?*

By design, scale-out provides:

- ▶ parallelism by automatically distributing load
- ▶ fault tolerance by rescheduling computation
- ▶ scalable storage with a distributed file system
- ▶ portability; Hadoop applications can run on any cluster
- ▶ availability because it can continually service clients
- ▶ scalability

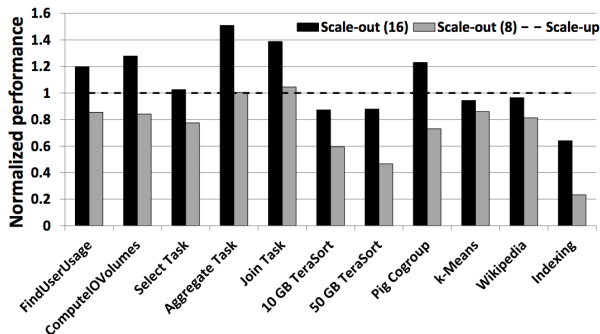
These properties may or may not affect performance...

...but they can't be ignored!

Scale-up vs. Scale-out Hadoop: Time to Rethink?

ACM Symposium on Cloud Computing '13 [2, 5]

- ▶ 10 “typical” jobs
- ▶ for today’s jobs, scale-up server > scale-out cluster



Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

Methodology


input




software



hardware



vs.



scale-up
vs.
scale-out

Michael Sevilla

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion



Methodology

input

- ▶ workload, input size

software

- ▶
- ▶
- ▶

hardware

- ▶


→ same workload, scale data

→


→

→

→



vs.



scale-up
vs.
scale-out

Michael Sevilla

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

Methodology

input

- ▶ workload, input size

software

- ▶ problem

▶

▶

hardware

▶

→ same workload, scale data


→ word count, sort

→


→

→

[3]



vs.



scale-up
vs.
scale-out

Michael Sevilla

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

Methodology

input

- ▶ workload, input size

software

- ▶ problem
- ▶ algorithm
- ▶

hardware

- ▶

→ same workload, scale data


→ word count, sort

→ methodology, functionality


→

[3]

→



vs.



scale-up
vs.
scale-out

Michael Sevilla

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

Methodology

input

- ▶ workload, input size

software

- ▶ problem
- ▶ algorithm
- ▶ scale-out properties

hardware

- ▶

→ same workload, scale data


→ word count, sort

[3]


→ methodology, functionality

→ implementations

→



vs.



scale-up
vs.
scale-out

Michael Sevilla

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

Methodology

input

- ▶ workload, input size

software

- ▶ problem
- ▶ algorithm
- ▶ scale-out properties

hardware

- ▶ processors, memory

→ same workload, scale data


→ word count, sort

[3]


→ methodology, functionality

→ implementations

→ \equiv compute contexts




vs.



Scale-up can Perform Better than Scale-out

scale-up
vs.
scale-out

Michael Sevilla



Introduction

- contributions
- challenges

Methodology

- parameters
- input
- software
- hardware

Analysis


- initial results
- implementations
- parallelism

Conclusion

Scale-up can Perform Better than Scale-out

scale-up
vs.
scale-out

Michael Sevilla



... but achieving scale-out properties changes the story!

- Other properties must be considered in comparison

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion


Achieving Scale-out Properties on Scale-up

scale-up
vs.
scale-out

Michael Sevilla

Phoenix: MapReduce runtime for multicore systems [4, 7, 6]

→ parallelism, port for methodology



Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

Distributed MultiThreaded Checkpointing (DMTCP) [1]

→ fault tolerance


Hadoop Distributed File System (HDFS)

→ scalable storage

Achieving Scale-out Properties on Scale-up

scale-up
vs.
scale-out

Michael Sevilla



Properties must be considered when comparing scale-out/up

Introduction

- contributions
- challenges

Methodology

- parameters
- input
- software
- hardware

Analysis


- initial results
- implementations**
- parallelism

Conclusion

Achieving Scale-out Properties on Scale-up

scale-up
vs.
scale-out

Michael Sevilla



Properties must be considered when comparing scale-out/up

- ▶ fault tolerance can make scale-up slower than scale-out

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis


initial results
implementations
parallelism

Conclusion

Achieving Scale-out Properties on Scale-up

scale-up
vs.
scale-out

Michael Sevilla



Properties must be considered when comparing scale-out/up

- ▶ fault tolerance can make scale-up slower than scale-out
- ▶ scalable storage may be the ultimate bottleneck

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis


initial results
implementations
parallelism

Conclusion

Achieving Scale-out Properties on Scale-up

scale-up
vs.
scale-out

Michael Sevilla



Properties must be considered when comparing scale-out/up

- ▶ fault tolerance can make scale-up slower than scale-out
- ▶ scalable storage may be the ultimate bottleneck
- ▶ parallelism limited by new job phases

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

Achieving Parallelism

scale-up
vs.
scale-out

Michael Sevilla

Introduction

contributions
challenges


Methodology

parameters
input
software
hardware

Analysis


initial results
implementations
parallelism

Conclusion



← Word count

Sort →



Achieving Parallelism

scale-up
vs.
scale-out

Michael Sevilla

Introduction

contributions
challenges


Methodology

parameters
input
software
hardware

Analysis


initial results
implementations
parallelism

Conclusion




← Word count

Sort →



Why is sort slower?!

scale-up
vs.
scale-out



- Sort is slower on scale-up because

1. new job phases
2. more key-value pairs

Conclusion

Compare scaling architectures (scale-up/out)

- ▶ comparison framework
 - encompasses {input, software, hardware} parameters
- ▶ achieving scale-out properties on scale-up

We show:

- ▶ must consider properties when comparing scale-out/up
- ▶ limitations of a scale-up computation framework

Questions?

scale-up
vs.
scale-out

Michael Sevilla

Special thanks to our anonymous reviewers for their helpful comments and suggestions.

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion

M. Sevilla, I. Nassi, K. Ioannidou, S. Brandt, C. Maltzahn. "A Framework for an In-depth Comparison of Scale-up and Scale-out". In Data-Intensive Scalable Computing Systems (DISCS), Denver, CO 2013.

scale-up
vs.
scale-out

Conclusion



Dmtcp: Transparent checkpointing for cluster computations and the desktop.

R. Appuswamy, C. Gkantsidis, D. Narayanan, O. Hodson, and A. Rowstron.

Scale-up vs scale-out for hadoop: Time to rethink?

In *Proceedings of the ACM Symposium on Cloud Computing*, 2013.

S. Huang, J. Huang, J. Dai, T. Xie, and B. Huang.

The HiBench Benchmark Suite: Characterization of the MapReduce-based Data Analysis.

In *ICDE Workshops*, pages 41–51, 2010.

C. Ranger, R. Raghuraman, A. Penmetsa, G. Bradski, and C. Kozyrakis.

Evaluating mapreduce for multi-core and multiprocessor systems.

In *Proceedings of the 2007 IEEE 13th International Symposium on High Performance Computer Architecture, HPCA '07*, pages 13–24, Washington, DC, USA, 2007. IEEE Computer Society.

A. Rowstron, D. Narayanan, A. Donnelly, G. O'Shea, and A. Douglas.

Nobody ever got fired for using hadoop on a cluster.

In *Proceedings of the 1st International Workshop on Hot Topics in Cloud Data Processing, HotCDP '12*, pages 2:1–2:5, New York, NY, USA, 2012. ACM.

J. Talbot, R. M. Yoo, and C. Kozyrakis.

Phoenix++: modular mapreduce for shared-memory systems.

In *Proceedings of the second international workshop on MapReduce and its applications*, MapReduce '11, pages 9–16, New York, NY, USA, 2011. ACM.

References II



R. M. Yoo, A. Romano, and C. Kozyrakis.

Phoenix rebirth: Scalable mapreduce on a large-scale shared-memory system.

In *Proceedings of the 2009 IEEE International Symposium on Workload Characterization (IISWC)*, IISWC '09, pages 198–207, Washington, DC, USA, 2009. IEEE Computer Society.

scale-up
vs.
scale-out

Michael Sevilla

Introduction

contributions
challenges

Methodology

parameters
input
software
hardware

Analysis

initial results
implementations
parallelism

Conclusion